


Introduction to Hadoop Distributed File System (HDFS)




Overview of HDFS Architecture

- *Able to explain HDFS design goals*
- *Factors impacting design*
- *Design approach*
- *Explain basic HDFS architecture*

HDFS Design Concept

- *Scalable distributed filesystem*
- *Distribute data on local disks on several nodes*
- *Low cost commodity hardware*



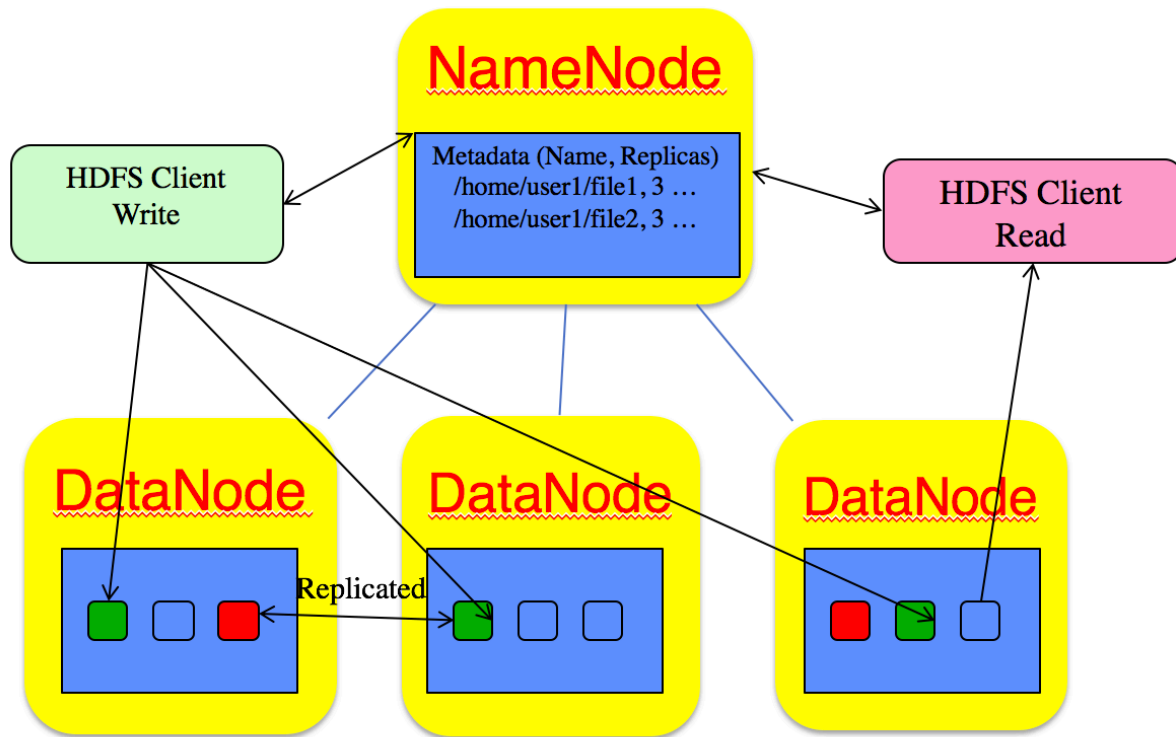
HDFS Design Factors

- *Hundreds/Thousands of nodes =>*
 - Need to handle node/disk failures
- *Portability across heterogeneous hardware/software*
- *Handle large data sets*
- *High throughput*

Approach to meet HDFS design goals

- *Simplified coherency model* – write once read many.
- *Data Replication* – helps handle hardware failures
- *Move computation close to data*
- *Relax POSIX requirements* – increase throughput

HDFS Architecture



Summary of HDFS Architecture


- **Single NameNode** - a master server that manages the file system namespace and regulates access to files by clients.
- **Multiple DataNodes** – typically one per node in the cluster. Functions:
 - Manage storage
 - Serving read/write requests from clients
 - Block creation, deletion, replication based on instructions from NameNode

Performance Envelope of HDFS

- *Able to determine number of blocks for a given file size*
- *Key HDFS and system components impacted by block size*
- *Impact of small files on HDFS and system*


Recall: HDFS Architecture

- *Distribute data on local disks on several nodes.*




HDFS block size

- *Default block size is 64MB*
- *Good for large files!*
- *So a 10GB file will be broken into:
 $10 \times 1024 / 64 = 160$ blocks.*



Importance of #blocks in a file

- **NameNode memory usage:** Every block represented as object (default replication this will be further increased 3X)
- **Number of map tasks:** data typically processed block at a time



Large #small files: Impact on NameNode

- **Memory usage** ~ 150 bytes per object
1 billion objects => 300GB memory!
- **Network load** – Number of checks with datanodes proportional to number of blocks.

Large #small files: Performance Impact

- **Map tasks** – depends on #blocks
10GB of data, 32k file size => 327680
map tasks
 - ⇒ lots of queued tasks
 - ⇒ large overhead of spin
up/tear down for each task
 - ⇒ Inefficient disk I/O with small
sizes

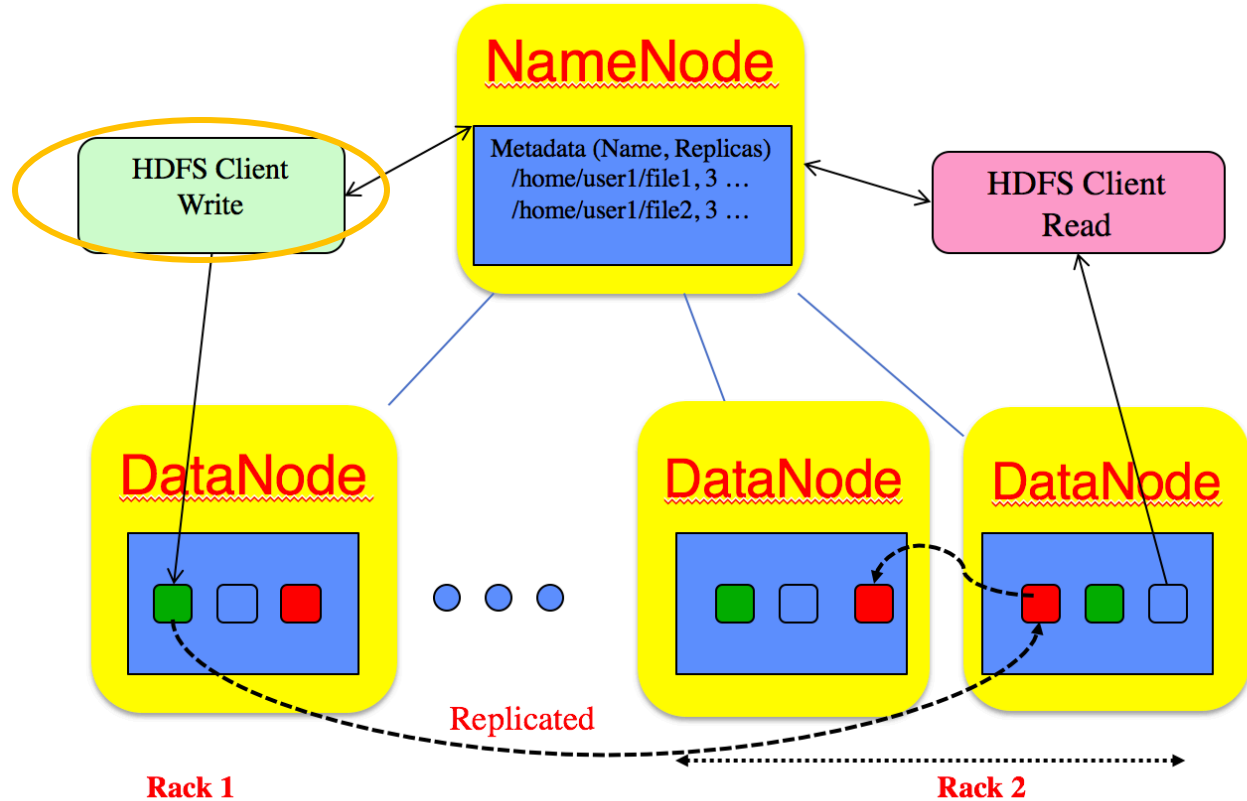
HDFS optimized for large files

- *Key takeaway – lots of small files is bad!*
- **Solutions:**
 - Merge/Concatenate files
 - Sequence files
 - HBase, HIVE configuration
 - CombineFileInputFormat

Write/Read processes on HDFS

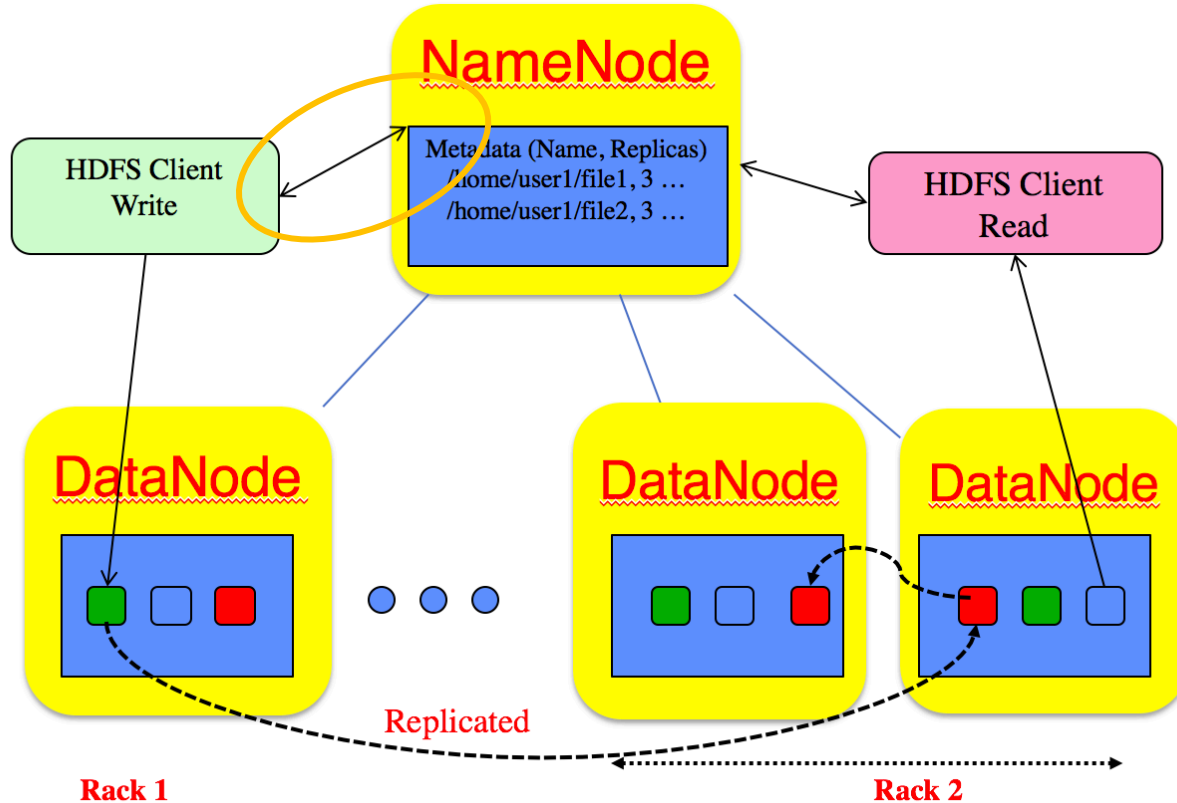
- *Able to explain write process in HDFS*
- *Detail the replication pipeline process*
- *Explain read process in HDFS*

Write Process in HDFS




*Client request
to create file*

Write Process in HDFS



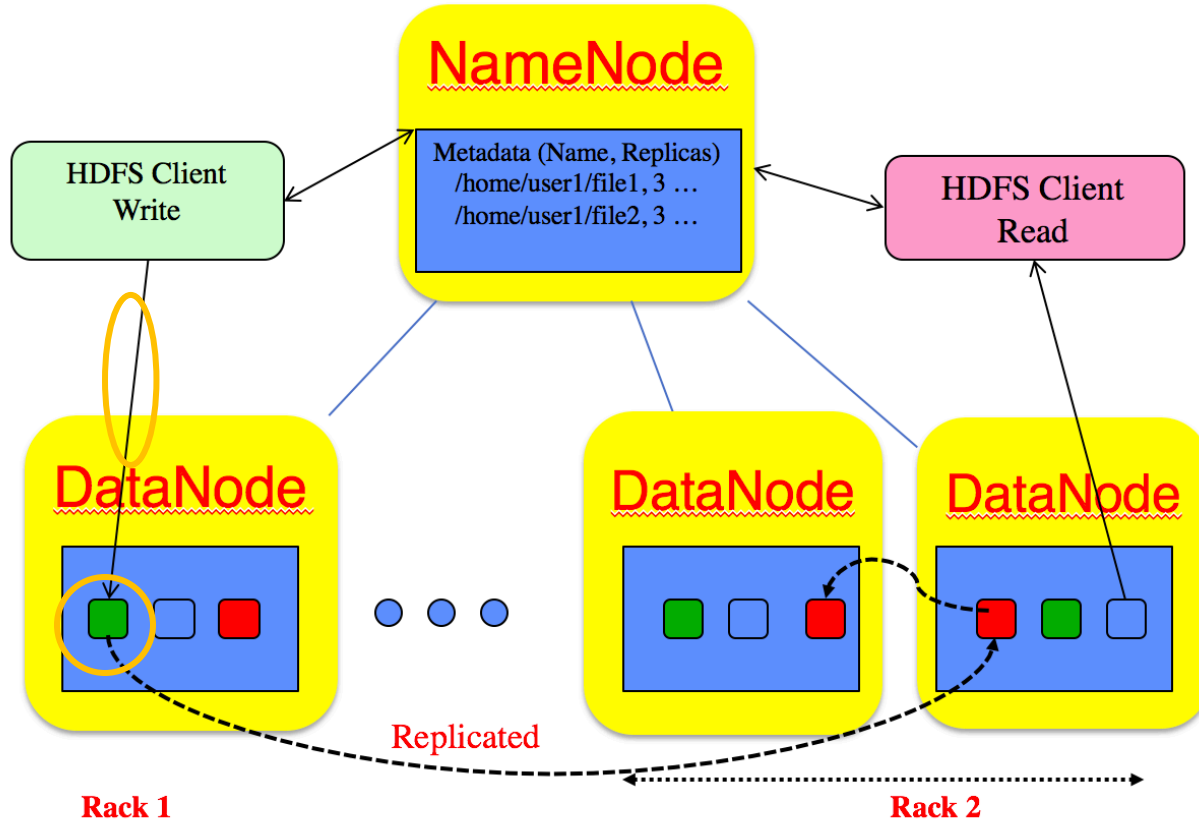
*NameNode
contacted
once a block
of data is
accumulated*

Write Process in HDFS



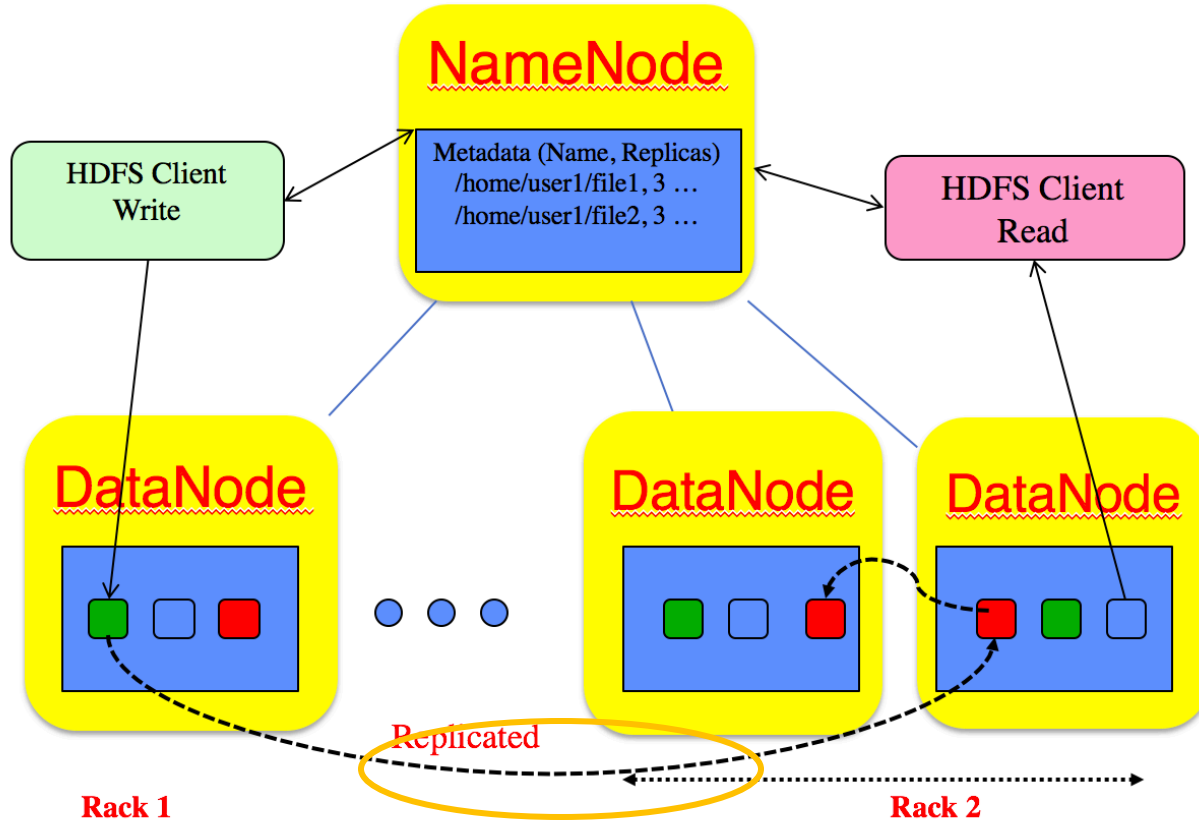
- *NameNode responds with list of DataNodes*
- *Rack aware*

Write Process in HDFS




- *First DataNode receives data, writes to local and forwards to second DataNode*

Write Process in HDFS



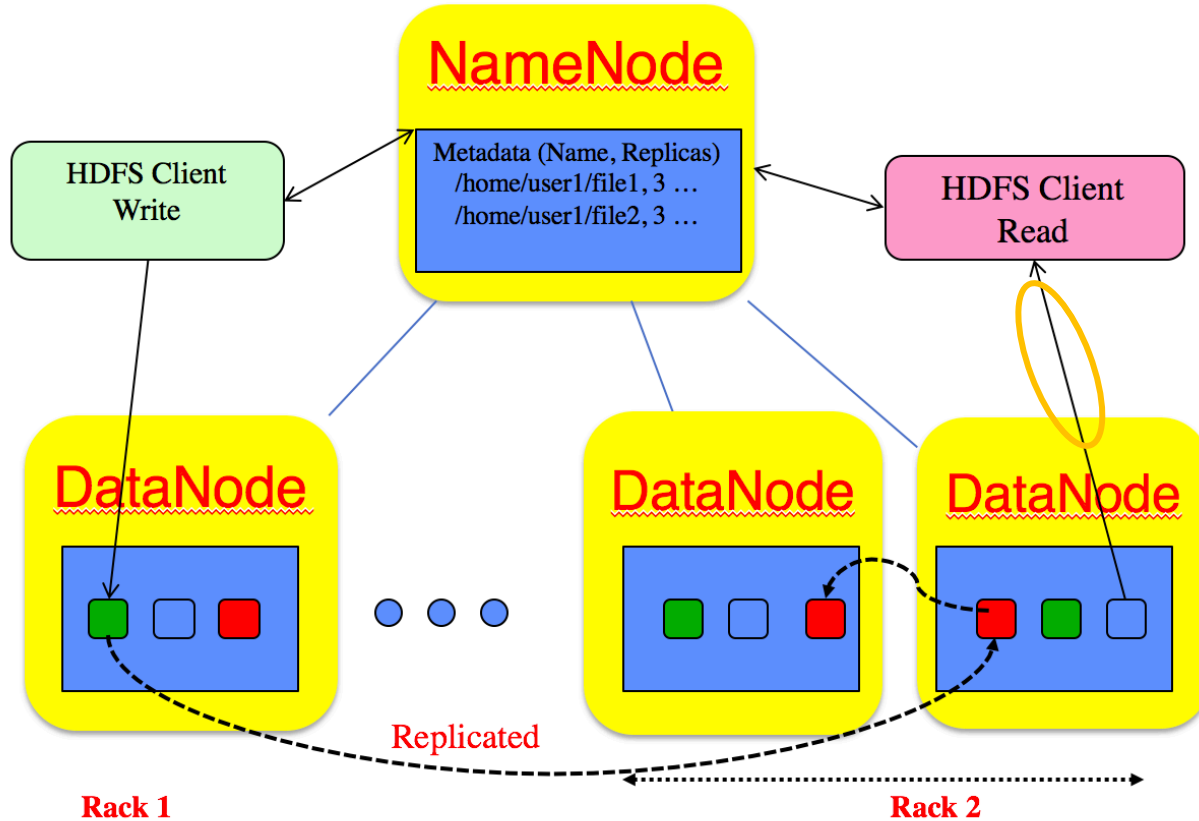
- First DataNode receives data, writes to local and forwards to second DataNode*

Write Process in HDFS



- *NameNode commits file creation into persistent store.*
- *Receives heartbeat and block reports*

Read Process in HDFS



- Client gets *DataNode list* from *NameNode*
- Read from *replica closest to reader.*